

TensorFlow 2.3.0 Support with Keras API

Overview

This document describes the modeling of networks using **TensorFlow 2 Keras API**. The ML models created and trained can be ported to and executed on MAX78000. Different types of Keras models with TensorFlow are supported, including high-level sequential, functional and sub-classing API. The following development approach has to be used:

1. Create Keras model using supported MAX78000 TF sub-classes which reflect hardware behavior and limit operations
2. Train model and store the model graph + weights into a **saved_model.pb** file
3. Use a Tensorflow-to-ONNX converter (**tf2onnx.convert**) to create an ONNX framework model from **saved_model.pb**
4. Quantize ONNX model weights and feed to MAX78000 synthesis tool to generate C code
5. Compile synthesized C code, load to MAX78000 and verify it

Setup

1- Install NVIDIA GPU drivers CUDA 10.1, CUDA Toolkit, CUPTI and cuDNN SDK 7.6 as described in Software requirements for Tensorflow:

<https://www.tensorflow.org/install/gpu>

2- Make sure that ~/.bash_profile includes path to CUDA and CUPTI:

```
export PATH="$PATH:/usr/local/cuda-10.1/bin"
export LD_LIBRARY_PATH="$LD_LIBRARY_PATH:/usr/local/cuda-10.1/extras/CUPTI/lib64"
```

3- To create a virtual environment please refer to section "Creating the Virtual Environment" of **[1]** document, including installation of ai8x-training and ai8x-synthesis requirements.

Supported MAX78000 TensorFlow Keras Subclasses:

ai85TF.py includes a set of customized TensorFlow 2 Keras subclasses to be used by any model that is designed to run on MAX78000.

Name	Description/Keras Equivalent
Conv1D	Generic Conv1D, padding_size=0
FusedConv1D	Conv1D with activation as None, padding_size=0
FusedConv1DReLU	Conv1D with activation as 'relu', padding_size=0
FusedMaxPoolConv1D	MaxPool1D, followed by Conv1D with activation as None, padding_size=0
FusedMaxPoolConv1DReLU	MaxPool1D, followed by Conv1D with activation as 'relu', padding_size=0
FusedAvgPoolConv1D	AveragePooling1D, followed by Conv1D with activation as None, padding_size=0
FusedAvgPoolConv1DReLU	AveragePooling1D followed by Conv1D with activation as 'relu', padding_size=0
MaxPool1D	MaxPool1D
AvgPool1D	AveragePooling1D
Conv2D	Generic Conv2D, padding_size=0
FusedConv2D	Conv2D with activation as None, padding_size=0
FusedConv2DReLU	Conv2D with activation as 'relu', padding_size=0
FusedMaxPoolConv2D	MaxPool2D, followed by Conv2D with activation as None, padding_size=0
FusedMaxPoolConv2DReLU	MaxPool2D, followed by Conv2D with activation as 'relu', padding_size=0
FusedAvgPoolConv2D	AveragePooling2D, followed by Conv2D with activation as None, padding_size=0
FusedAvgPoolConv2DReLU	AveragePooling2D followed by Conv2D with activation as 'relu', padding_size=0
MaxPool2D	MaxPool2D
AvgPool2D	AveragePooling2D
Dense	Generic Dense
FusedDense	Dense with activation as None
FusedDenseReLU	Dense with activation as 'relu'

Limitations of supported operations:

Conv2D:

- Kernel sizes must be 1×1 or 3×3.
- Padding can be 0, 1, or 2 (default: padding_size = 0).
- Stride is fixed to 1. Pooling, including 1×1, can be used to achieve a stride other than 1.

Conv1D:

- Kernel sizes must be 1 through 9.
- Padding can be 0, 1, or 2 (default: padding_size = 0).
- Stride is fixed to 1. Pooling, including 1, can be used to achieve a stride other than 1.

Conv2DTranspose:

- Kernel sizes must be 3×3.
- Padding can be 0, 1, or 2 (default: padding_size = 0).
- Stride is fixed to 2

Pooling:

- Both max pooling and average pooling are available, with or without convolution.
- Pooling does not support padding.
- Pooling strides can be 1 through 16. For 2D pooling, the stride is the same for both dimensions.
- For 2D pooling, supported pooling kernel sizes are 1×1 through 16×16, including non-square kernels. 1D pooling supports kernels from 1 through 16. *Note: 1×1 kernels can be used when a convolution stride other than 1 is desired.*
- The number of input channels must not exceed 1024.
- The number of output channels must not exceed 1024.

For more details of MAX7800 HW related limitations please check **[1]** document.

Training Models

Following bash scripts are provided to download the dataset, train the model and to convert to onnx format:

```
train_cifar10.sh
train_kws20.sh
train_mnist.sh
train_cifar100.sh
train_fashionmnist.sh
train_rock.sh
```

Example: (train_mnist.sh)

```
python train.py --epochs 100 --batch_size 256 --optimizer Adam --lr 0.001 --
model mnist_model --dataset mnist --save-sample 1
python -m tf2onnx.convert --saved-model export/mnist --opset 10 --output
export/mnist/saved_model.onnx
```

The script automatically downloads corresponding dataset and process and copy into **/data/** if needed and starts training. Training progress and results including checkpoint and a sample prediction for one test data in HWC format will be stored in log file inside **/logs/** directory. The model graph and weights are stored as **saved_model.pb** file in log directory, as well as in **/export/** directory.

Command-line args for Training

User can modify following training command-line parameters:

Training parameter	Description
--epochs	Number of training epochs (default: 100)
--batch_size	Training batch size (default: 32)
--optimizer	Optimizer type: Adam or SGD (default: Adam)
--lr	Initial learning rate (default: 0.0001). During training learning rate is adjusted according to schedule in model
--model	Model name
--dataset	Dataset name
--save-sample	Save input sample with specified index in .npy format in /export/ folder for verification in synthesis in
--save-sample-per-class	Save one input sample for each class in .npy format in logs folder to be used for verification
--channel-first	Save sample in channel first format in multi-channel cases (default: channel last as native form on Tensorflow), suitable to be used by synthesis script (NCHW)
--swap	if --channel-first is selected, this option swaps order of H,W in saved sample data. This is only needed if first layer is a conv1d. (default: no swap)
--metrics	Metrics used in compiling model (default: accuracy)

Once training is complete, the model is converted to ONNX format and stored in **/export/**.

Models

Model examples are located in **models** directory. Each model includes a Tensorflow Keras sequential model, as well as the callback function for learning rate adjustment scheduler:

```
models/cifar10_model.py
models/cifar100_model.py
models/fashionmnist_model.py
models/kws20_model.py
models/mnist_model.py
models/rock_model.py
```

Learning rate scheduler

Each model script includes the lr scheduler callback function to be used for that model. The default schedule is ReduceLROnPlateau with following parameters:

```
lr_schedule = tf.keras.callbacks.ReduceLROnPlateau(
    monitor='val_accuracy',
    mode='max',
    factor=0.2,
    patience=3,
    verbose=1,
    min_lr=1e-5)
```

Datasets

Dataset scripts are located in **/dataset/** directory and used to download the dataset and create a processed **.npz** dataset file (if needed) in **/data/** to be used by training script:

```
datasets/cifar10.py
datasets/cifar100.py
datasets/fashionmnist.py
datasets/kws20.py
datasets/mnist.py
datasets/rock.py
```

Datasets include training, validation and test images and labels. Images are in [-128,127] range when created by dataset scripts.

In training script, they are normalized to [-0.5,0.5] and fed to the network. However, in synthesis script, samples are expected in [-128,127] range and normalized to [-1,+1] range to mimic hardware. To accommodate x2 scaling, the first layer weights are scaled by 1/2 by default, unless **--keep-first** is selected in synthesis.

Examples

MNIST model

The MNIST model is an example of Keras model with sequential API and it recognizes 28x28 images of handwritten digits from 0 to 9.

```
model = tf.keras.models.Sequential([
    tf.keras.Input(shape=(28, 28)),
    tf.keras.layers.Reshape(target_shape=(28, 28, 1)),
    ai8xTF.FusedConv2DReLU(
        filters=60,
```

```

        kernel_size=3,
        strides=1,
        padding_size=1,
        use_bias=False),
ai8xTF.FusedMaxPoolConv2DReLU(
    filters=60,
    kernel_size=3,
    strides=1,
    padding_size=2,
    pool_size=2,
    pool_strides=2,
    use_bias=False),
ai8xTF.FusedMaxPoolConv2DReLU(
    filters=56,
    kernel_size=3,
    strides=1,
    padding_size=1,
    pool_size=2,
    pool_strides=2,
    use_bias=False),
ai8xTF.FusedAvgPoolConv2DReLU(
    filters=12,
    kernel_size=3,
    strides=1,
    padding_size=1,
    pool_size=2,
    pool_strides=2,
    use_bias=False),
tf.keras.layers.Flatten(),
ai8xTF.FusedDense(10, wide=True, use_bias=True),
])

```

To train MNIST model execute following script:

```
$ bash train_mnist.sh
```

Training progress, accuracy results and confusion table are reported and stored in log file.

```

Epoch 98/100
211/211 - 2s - loss: 0.0022 - accuracy: 1.0000 - val_loss: 0.0279 -
val_accuracy: 0.9908
Epoch 99/100
211/211 - 2s - loss: 0.0022 - accuracy: 1.0000 - val_loss: 0.0279 -
val_accuracy: 0.9908
Epoch 100/100
211/211 - 2s - loss: 0.0021 - accuracy: 1.0000 - val_loss: 0.0278 -
val_accuracy: 0.9908
188/188 - 0s - loss: 0.0302 - accuracy: 0.9900
Test Accuracy: 0.9900000095367432
Confusion Matrix:
tf.Tensor(
[[619  0  1  1  0  0  1  0  2  0]
 [ 0 654  0  0  0  0  0  0  0  0]
 [ 0  1 569  0  0  0  0  1  0  1]
 [ 0  0  2 584  0  2  0  0  1  0]
 [ 1  0  0  0 572  0  0  1  0  6]

```

```
[ 1  0  0  1  0 544  2  0  3  0]
[ 2  0  1  1  1  2 572  0  1  0]
[ 0  1  3  0  1  0  0 627  0  1]
[ 1  0  2  1  1  0  0  0 580  0]
[ 0  0  0  1  7  1  0  3  1 619]], shape=(10, 10), dtype=int32)
```

At end of training a summary of the model is reported and model graph/weights are stored as **saved_model.pb** file

Model: "sequential"

Layer (type)	Output Shape	Param #
reshape (Reshape)	(None, 28, 28, 1)	0
fused_conv2d_re_lu (FusedCon	(None, 28, 28, 60)	540
fused_max_pool_conv2d_re_lu	(None, 16, 16, 60)	32400
fused_max_pool_conv2d_re_lu_	(None, 8, 8, 56)	30240
fused_avg_pool_conv2d_re_lu	(None, 4, 4, 12)	6048
flatten (Flatten)	(None, 192)	0
fused_dense (FusedDense)	(None, 10)	1930
Total params: 71,158		
Trainable params: 71,158		
Non-trainable params: 0		

Note: Empty class may be included as part of subclasses in the sequential model. However, it is not needed and skipped in serialization.

Additionally, the model is converted to ONNX format:

```
export/mnist/saved_model.pb
export/mnist/saved_model.onnx
```

Fashion MNIST model

This model demonstrates recognition of 10 28x28 fashion images: **T-shirt/top, trouser, pullover, dress, coat, sandal, shirt, sneaker, bag, ankle boot.**

The Fashion MNIST model is an example of Keras model with functional API:

```
# create a functional model
input_layer = tf.keras.Input(shape=(28, 28))

reshape = tf.keras.layers.Reshape(target_shape=(28, 28, 1))(input_layer)
conv1 = ai8xTF.FusedConv2DReLU(
```

```

        filters=60,
        kernel_size=3,
        strides=1,
        padding_size=1)(reshape)

conv2 = ai8xTF.FusedMaxPoolConv2DReLU(
    filters=60,
    kernel_size=3,
    strides=1,
    padding_size=2,
    pool_size=2,
    pool_strides=2)(conv1)

# dropout1= tf.keras.layers.Dropout(0.2)(conv2)

conv3 = ai8xTF.FusedMaxPoolConv2DReLU(
    filters=56,
    kernel_size=3,
    strides=1,
    padding_size=1,
    pool_size=2,
    pool_strides=2)(conv2)

conv4 = ai8xTF.FusedAvgPoolConv2DReLU(
    filters=12,
    kernel_size=3,
    strides=1,
    padding_size=1,
    pool_size=2,
    pool_strides=2)(conv3)

flat = tf.keras.layers.Flatten(input_shape=(28, 28))(conv4)

output_layer = ai8xTF.FusedDense(10, wide=True)(flat)

model = tf.keras.Model(inputs=[input_layer], outputs=[output_layer])

```

To train fashion MNIST model execute following script:

```
$ bash train_fashionmnist.sh
```

Training progress, accuracy results and confusion table are reported and stored in log file.

```

Epoch 97/100
211/211 - 2s - loss: 0.1520 - accuracy: 0.9490 - val_loss: 0.2439 -
val_accuracy: 0.9121
Epoch 98/100
211/211 - 2s - loss: 0.1517 - accuracy: 0.9492 - val_loss: 0.2441 -
val_accuracy: 0.9113
Epoch 99/100
211/211 - 2s - loss: 0.1515 - accuracy: 0.9491 - val_loss: 0.2440 -
val_accuracy: 0.9108
Epoch 100/100
211/211 - 2s - loss: 0.1512 - accuracy: 0.9493 - val_loss: 0.2441 -
val_accuracy: 0.9107
188/188 - 0s - loss: 0.2329 - accuracy: 0.9148

```



```

Test Accuracy: 0.9148333072662354
Confusion Matrix:
tf.Tensor(
[[522   1  11  14   1   0  46   0   2   0]
 [  0 597   0  10   0   0   1   0   0   0]
 [  9   0 545   3  31   0  22   0   1   0]
 [ 15   3   2 536  17   0  13   0   1   0]
 [  0   0  24  22 546   0  33   0   2   0]
 [  0   0   0   0   0 601   1  11   4   4]
 [ 60   2  32  17  35   0 467   0   6   0]
 [  0   0   0   0   0   7   0 533   0  10]
 [  2   1   2   0   2   2   3   0 578   0]
 [  0   0   0   0   0   3   0  22   1 564]], shape=(10, 10), dtype=int32)

```

At end of training a summary of the model is reported and model graph/weights are stored as **saved_model.pb** file

Model: "functional_1"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 28, 28)]	0
reshape (Reshape)	(None, 28, 28, 1)	0
fused_conv2d_re_lu (FusedCon	(None, 28, 28, 60)	600
fused_max_pool_conv2d_re_lu	(None, 16, 16, 60)	32460
fused_max_pool_conv2d_re_lu_	(None, 8, 8, 56)	30296
fused_avg_pool_conv2d_re_lu	(None, 4, 4, 12)	6060
flatten (Flatten)	(None, 192)	0
fused_dense (FusedDense)	(None, 10)	1930
=====		
Total params: 71,346		
Trainable params: 71,346		
Non-trainable params: 0		

Note: Empty class may be included as part of subclasses in the sequential model. However, it is not needed and skipped in serialization.

Additionally, the model is converted to ONNX format

```

export/fashionmnist/saved_model.pb
export/fashionmnist/saved_model.onnx

```

CIFAR10 model

The CIFAR-10 dataset consists of 60000 32x32 color images in 10 classes: **plane, car, bird, cat, deer, dog, frog, horse, ship, truck**.

<https://www.cs.toronto.edu/~kriz/cifar.html>

The CIFAR10 model is an example of Keras model with sequential API:

```
model = tf.keras.models.Sequential([
    tf.keras.Input(shape=(32, 32, 3)),
    ai8xTF.FusedConv2DReLU(
        filters=60, kernel_size=3, strides=1, padding_size=1, use_bias=False),
    ai8xTF.FusedMaxPoolConv2DReLU(
        filters=60,
        kernel_size=3,
        strides=1,
        padding_size=1,
        pool_size=2,
        pool_strides=2,
        use_bias=False),
    ai8xTF.FusedMaxPoolConv2DReLU(
        filters=56,
        kernel_size=3,
        strides=1,
        padding_size=1,
        pool_size=2,
        pool_strides=2,
        use_bias=False),
    ai8xTF.FusedAvgPoolConv2DReLU(
        filters=12,
        kernel_size=3,
        strides=1,
        padding_size=1,
        pool_size=2,
        pool_strides=2,
        use_bias=False),
    tf.keras.layers.Flatten(),
    ai8xTF.FusedDense(10, wide=True, use_bias=False),
])
```

To train CIFAR10 model execute following script:

```
$ bash train_cifar10.sh
```

Training progress, accuracy results and confusion table are reported and stored in log file.

```
Epoch 98/100
704/704 - 5s - loss: 0.5455 - accuracy: 0.8174 - val_loss: 0.8520 -
val_accuracy: 0.7002
Epoch 99/100
704/704 - 6s - loss: 0.5443 - accuracy: 0.8179 - val_loss: 0.8533 -
val_accuracy: 0.7017
Epoch 100/100
704/704 - 6s - loss: 0.5436 - accuracy: 0.8173 - val_loss: 0.8515 -
val_accuracy: 0.7040
157/157 - 0s - loss: 0.8420 - accuracy: 0.7036
Test Accuracy: 0.7035999894142151
Confusion Matrix:
tf.Tensor(
[[356 15 21 10 6 3 3 8 31 23]
```

```
[ 13 402  5  2  1  1  3  4 17 39]
[ 41  3 311 32 43 36 36 13  6  6]
[ 13  1 35 243 37 126 37 21  5  5]
[ 16  1 22 20 335 17 26 41  4  4]
[  3  1 27 91 26 312  9 31  1  3]
[  7  2 31 26 33 12 367  4  1  2]
[  9  2 20 15 36 30  3 382  3  9]
[ 41 20  2  6  1  0  4  3 416 10]
[ 15 47  4  6  7  5  2  5 15 394]], shape=(10, 10), dtype=int32)
```

At end of training a summary of the model is reported and model graph/weights are stored as **saved_model.pb** file

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
fused_conv2d_re_lu (FusedConv2DReLU)	(None, 32, 32, 60)	1620
fused_max_pool_conv2d_re_lu (FusedMaxPool2DReLU)	(None, 16, 16, 60)	32400
fused_max_pool_conv2d_re_lu (FusedMaxPool2DReLU)	(None, 8, 8, 56)	30240
fused_avg_pool_conv2d_re_lu (FusedAvgPool2DReLU)	(None, 4, 4, 12)	6048
flatten (Flatten)	(None, 192)	0
fused_dense (FusedDense)	(None, 10)	1920
=====		
Total params: 72,228		
Trainable params: 72,228		
Non-trainable params: 0		

Note: Empty class may be included as part of subclasses in the sequential model. However, it is not needed and skipped in serialization.

Additionally, the model is converted to ONNX format

```
export/cifar10/saved_model.pb
export/cifar10/saved_model.onnx
```

CIFAR100 model

The CIFAR100 model classifies 100 32x32 color images from 60K dataset.

<https://www.cs.toronto.edu/~kriz/cifar.html>

The CIFAR100 model is an example of Keras model with sequential API:

```
model = tf.keras.models.Sequential([
    tf.keras.Input(shape=(32, 32, 3)),
    ai8xTF.FusedConv2DReLU(
        filters=60, kernel_size=3, strides=1, padding_size=1, use_bias=False),
```

```

ai8xTF.FusedMaxPoolConv2DReLU(
    filters=60,
    kernel_size=3,
    strides=1,
    padding_size=1,
    pool_size=2,
    pool_strides=2,
    use_bias=False),
ai8xTF.FusedMaxPoolConv2DReLU(
    filters=56,
    kernel_size=3,
    strides=1,
    padding_size=1,
    pool_size=2,
    pool_strides=2,
    use_bias=False),
ai8xTF.FusedAvgPoolConv2DReLU(
    filters=12,
    kernel_size=3,
    strides=1,
    padding_size=1,
    pool_size=2,
    pool_strides=2,
    use_bias=False),
tf.keras.layers.Flatten(),
ai8xTF.FusedDense(10, wide=True, use_bias=False),
])

```

To train CIFAR100 model execute following script:

```
$ bash train_cifar100.sh
```

Training progress, accuracy results and confusion table are reported and stored in log file.

```

Epoch 99/100
1407/1407 - 11s - loss: 1.1360 - accuracy: 0.7284 - val_loss: 2.8744 -
val_accuracy: 0.3305
Epoch 100/100
1407/1407 - 11s - loss: 1.1316 - accuracy: 0.7310 - val_loss: 2.8761 -
val_accuracy: 0.3283
157/157 - 0s - loss: 2.8540 - accuracy: 0.3338
Test Accuracy: 0.33379998803138733
Confusion Matrix:
tf.Tensor(
[[34  2  0 ...  0  0  0]
 [ 0 18  1 ...  0  0  0]
 [ 0  0 19 ...  1  4  0]
 ...
 [ 0  0  0 ... 11  0  0]
 [ 0  0  2 ...  1 18  0]
 [ 0  0  0 ...  1  0  8]], shape=(100, 100), dtype=int32)

```

At end of training a summary of the model is reported and model graph/weights are stored as **saved_model.pb** file

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
=====		
fused_conv2d_re_lu (FusedCon	(None, 32, 32, 16)	432
fused_conv2d_re_lu_1 (FusedC	(None, 32, 32, 20)	2880
fused_conv2d_re_lu_2 (FusedC	(None, 32, 32, 20)	3600
fused_conv2d_re_lu_3 (FusedC	(None, 32, 32, 20)	3600
fused_max_pool_conv2d_re_lu	(None, 16, 16, 20)	3600
fused_conv2d_re_lu_4 (FusedC	(None, 16, 16, 20)	3600
fused_conv2d_re_lu_5 (FusedC	(None, 16, 16, 44)	7920
fused_max_pool_conv2d_re_lu_	(None, 8, 8, 48)	19008
fused_conv2d_re_lu_6 (FusedC	(None, 8, 8, 48)	20736
fused_max_pool_conv2d_re_lu_	(None, 4, 4, 96)	41472
fused_max_pool_conv2d_re_lu_	(None, 2, 2, 512)	49152
fused_conv2d_re_lu_7 (FusedC	(None, 2, 2, 128)	65536
fused_max_pool_conv2d_re_lu_	(None, 1, 1, 128)	147456
conv2d_13 (Conv2D)	(None, 1, 1, 100)	12800
flatten (Flatten)	(None, 100)	0
=====		
Total params: 381,792		
Trainable params: 381,792		
Non-trainable params: 0		

Note: Empty class may be included as part of subclasses in the sequential model. However, it is not needed and skipped in serialization.

Additionally, the model is converted to ONNX format

```
export/cifar100/saved_model.pb
export/cifar100/saved_model.onnx
```

KWS20 model

The KWS20 model uses the 2nd version of Google speech commands dataset which consists of 35 keywords and more than 100K utterances.

https://storage.cloud.google.com/download.tensorflow.org/data/speech_commands_v0.02.tar.gz

This model demonstrates recognition of 20 keywords: 'up', 'down', 'left', 'right', 'stop', 'go', 'yes', 'no', 'on', 'off', 'one', 'two', 'three', 'four', 'five', 'six', 'seven', 'eight', 'nine', 'zero'. The rest of keywords are placed into category as "unknown".

The KWS20 model is an example of Keras model with sequential API:

```
model = tf.keras.models.Sequential([
    # Need to specify the input shape if you want to show it in model summary
    tf.keras.Input(shape=(128, 128)),
    ai8xTF.FusedConv1DReLU(
        filters=100,
        kernel_size=1,
        strides=1,
        padding_size=0,
        kernel_regularizer=regularizer,
        activity_regularizer=activity_regularizer,
        use_bias=False),
    ai8xTF.FusedConv1DReLU(
        filters=100,
        kernel_size=1,
        strides=1,
        padding_size=0,
        kernel_regularizer=regularizer,
        activity_regularizer=activity_regularizer,
        use_bias=False),
    ai8xTF.FusedConv1DReLU(
        filters=50,
        kernel_size=1,
        strides=1,
        padding_size=0,
        kernel_regularizer=regularizer,
        activity_regularizer=activity_regularizer,
        use_bias=False),
    ai8xTF.FusedConv1DReLU(
        filters=16,
        kernel_size=1,
        strides=1,
        padding_size=0,
        kernel_regularizer=regularizer,
        activity_regularizer=activity_regularizer,
        use_bias=False),

    # Conversion 1D to 2D
    tf.keras.layers.Reshape(target_shape=(8, 16, 16)),
    ai8xTF.FusedConv2DReLU(
        filters=32,
        kernel_size=3,
        strides=1,
        padding_size=1,
        kernel_regularizer=regularizer,
        activity_regularizer=activity_regularizer,
        use_bias=False),
    ai8xTF.FusedConv2DReLU(
        filters=64,
        kernel_size=3,
        strides=1,
        padding_size=1,
```

```

        kernel_regularizer=regularizer,
        activity_regularizer=activity_regularizer,
        use_bias=False),
ai8xTF.FusedConv2DReLU(
    filters=64,
    kernel_size=3,
    strides=1,
    padding_size=1,
    kernel_regularizer=regularizer,
    activity_regularizer=activity_regularizer,
    use_bias=False),
ai8xTF.FusedConv2DReLU(
    filters=30,
    kernel_size=3,
    strides=1,
    padding_size=1,
    kernel_regularizer=regularizer,
    activity_regularizer=activity_regularizer,
    use_bias=False),
ai8xTF.FusedConv2DReLU(
    filters=7,
    kernel_size=3,
    strides=1,
    padding_size=1,
    kernel_regularizer=regularizer,
    activity_regularizer=activity_regularizer,
    use_bias=False),
tf.keras.layers.Flatten(),
ai8xTF.FusedDense(
    21, wide=True,
    use_bias=False,
    kernel_regularizer=regularizer,
    activity_regularizer=activity_regularizer),
])

```

To train KWS20 model execute following script:

```
$ bash train_kws20.sh
```

Training progress, accuracy results and confusion table are reported and stored in log file.

```

Epoch 198/200
667/667 - 5s - loss: 0.3848 - accuracy: 0.9528 - val_loss: 0.6943 -
val_accuracy: 0.8573
Epoch 199/200
667/667 - 5s - loss: 0.3847 - accuracy: 0.9522 - val_loss: 0.7012 -
val_accuracy: 0.8545
Epoch 200/200
667/667 - 4s - loss: 0.3842 - accuracy: 0.9527 - val_loss: 0.6977 -
val_accuracy: 0.8561
319/319 - 1s - loss: 0.7320 - accuracy: 0.8514
Test Accuracy: 0.8514375686645508
Confusion Matrix:
tf.Tensor(
[[ 323   0   0   0   6   4   0   2   4  15   0   1   0   0   0
   0   0   1   0   0  17]

```

```

[ 2 303 0 0 1 13 0 11 1 0 1 2 0 1 0
1 1 0 8 0 34]
[ 1 0 304 8 0 0 13 3 0 2 1 0 1 0 0
0 0 1 3 1 23]
[ 1 1 7 297 0 0 0 0 2 0 5 0 2 1 6
0 0 1 14 0 21]
[ 4 1 0 0 350 2 0 0 0 2 0 1 0 2 0
3 8 0 0 0 9]
[ 4 11 0 0 3 273 0 17 1 1 0 6 0 5 0
0 0 3 1 0 46]
[ 0 2 8 0 0 0 364 2 0 3 2 0 1 0 0
3 0 1 0 2 15]
[ 2 12 0 0 0 21 2 307 0 0 2 1 0 1 0
0 0 0 7 2 30]
[ 2 4 0 0 0 2 0 0 298 12 9 1 1 2 15
1 0 0 2 0 12]
[ 24 0 1 0 0 2 2 0 9 284 1 1 0 4 5
0 0 0 0 0 7]
[ 1 1 2 4 0 2 1 3 7 0 320 0 0 1 2
0 0 0 7 0 33]
[ 1 1 1 0 1 7 1 0 0 0 0 319 3 4 0
2 3 0 1 11 18]
[ 0 0 0 3 0 5 1 0 1 0 0 9 289 0 0
4 3 6 1 3 32]
[ 0 0 0 0 0 6 0 0 6 2 3 0 1 286 1
3 0 0 0 3 48]
[ 1 2 2 10 1 0 0 0 9 2 2 0 3 0 321
2 1 1 3 0 26]
[ 0 0 0 0 0 0 2 0 0 0 0 1 1 0 0
349 0 3 0 1 7]
[ 0 6 0 0 5 0 0 0 2 0 0 6 2 1 1
7 332 0 1 2 30]
[ 1 0 0 0 0 2 2 0 0 0 0 2 6 0 1
5 0 349 0 0 15]
[ 0 2 5 9 0 3 0 7 3 0 6 0 2 0 3
0 0 0 319 1 28]
[ 0 1 1 0 0 1 0 1 0 0 0 7 4 1 0
3 4 0 0 358 18]
[ 20 32 16 12 9 32 10 21 15 14 17 10 36 37 17
5 10 7 18 18 2332]], shape=(21, 21), dtype=int32)

```

At end of training a summary of the model is reported and model graph/weights are stored as **saved_model.pb** file

Model: "sequential"		
Layer (type)	Output Shape	Param #
=====		
fused_conv1d_re_lu (FusedCon	(None, 128, 100)	12800
fused_conv1d_re_lu_1 (FusedC	(None, 128, 100)	10000
fused_conv1d_re_lu_2 (FusedC	(None, 128, 50)	5000
fused_conv1d_re_lu_3 (FusedC	(None, 128, 16)	800
reshape (Reshape)	(None, 8, 16, 16)	0

fused_conv2d_re_lu	(FusedCon (None, 8, 16, 32)	4608
fused_conv2d_re_lu_1	(FusedC (None, 8, 16, 64)	18432
fused_conv2d_re_lu_2	(FusedC (None, 8, 16, 64)	36864
fused_conv2d_re_lu_3	(FusedC (None, 8, 16, 30)	17280
fused_conv2d_re_lu_4	(FusedC (None, 8, 16, 7)	1890
flatten	(Flatten) (None, 896)	0
fused_dense	(FusedDense) (None, 21)	18816
=====		
Total params: 126,490		
Trainable params: 126,490		
Non-trainable params: 0		

Note: Empty class may be included as part of subclasses in the sequential model. However, it is not needed and skipped in serialization.

Additionally, the model is converted to ONNX format

```
export/kws20/saved_model.pb
export/kws20/saved_model.onnx
```

Rock-Paper-Scissor model

This model demonstrates recognition of images of hands playing rock, paper, scissor popular game.

https://www.tensorflow.org/datasets/catalog/rock_paper_scissors

The RPS model is an example of Keras model with sequential API:

```
IMG_SIZE = 64 # All images will be resized to 120x120
# Setup model
model = tf.keras.models.Sequential([
    tf.keras.Input(shape=(IMG_SIZE, IMG_SIZE, 3)),
    ai8xTF.FusedConv2DReLU(
        filters=15,
        kernel_size=3,
        strides=1,
        padding_size=1
    ),
    ai8xTF.FusedMaxPoolConv2DReLU(
        filters=30,
        kernel_size=3,
        strides=1,
        padding_size=1,
        pool_size=2,
        pool_strides=2
    ),
```

```

tf.keras.layers.Dropout(0.2),
ai8xTF.FusedMaxPoolConv2DReLU(
    filters=60,
    kernel_size=3,
    strides=1,
    padding_size=1,
    pool_size=2,
    pool_strides=2
),
ai8xTF.FusedMaxPoolConv2DReLU(
    filters=30,
    kernel_size=3,
    strides=1,
    padding_size=1,
    pool_size=2,
    pool_strides=2
),
ai8xTF.FusedMaxPoolConv2DReLU(
    filters=30,
    kernel_size=3,
    strides=1,
    padding_size=1,
    pool_size=2,
    pool_strides=2
),
ai8xTF.FusedConv2DReLU(
    filters=30,
    kernel_size=3,
    strides=1,
    padding_size=1
),
tf.keras.layers.Flatten(),
tf.keras.layers.Dropout(0.2),
ai8xTF.FusedDense(3, wide=True)
])

```

To train RPS model execute following script:

```
$ bash train_rock.sh
```

Training progress, accuracy results and confusion table are reported and stored in log file.

```

Epoch 98/100
79/79 - 0s - loss: 6.1023e-05 - accuracy: 1.0000 - val_loss: 0.3391 -
val_accuracy: 0.9140
Epoch 99/100
79/79 - 1s - loss: 5.7724e-05 - accuracy: 1.0000 - val_loss: 0.3437 -
val_accuracy: 0.9140
Epoch 100/100
79/79 - 0s - loss: 6.5221e-05 - accuracy: 1.0000 - val_loss: 0.3369 -
val_accuracy: 0.9140
6/6 - 0s - loss: 0.3191 - accuracy: 0.9140
Test Accuracy: 0.9139785170555115
Confusion Matrix:
tf.Tensor(
[[61  0  0]
 [ 8 47  8]
 [ 0  0 62]], shape=(3, 3), dtype=int32)

```

At end of training a summary of the model is reported and model graph/weights are stored as **saved_model.pb** file

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
fused_conv2d_re_lu (FusedCon	(None, 64, 64, 15)	420
fused_max_pool_conv2d_re_lu	(None, 32, 32, 30)	4080
dropout (Dropout)	(None, 32, 32, 30)	0
fused_max_pool_conv2d_re_lu_	(None, 16, 16, 60)	16260
fused_max_pool_conv2d_re_lu_	(None, 8, 8, 30)	16230
fused_max_pool_conv2d_re_lu_	(None, 4, 4, 30)	8130
fused_conv2d_re_lu_1 (FusedC	(None, 4, 4, 30)	8130
flatten (Flatten)	(None, 480)	0
dropout_1 (Dropout)	(None, 480)	0
fused_dense (FusedDense)	(None, 3)	1443
=====		
Total params: 54,693		
Trainable params: 54,693		
Non-trainable params: 0		

Note: Empty class may be included as part of subclasses in the sequential model. However, it is not needed and skipped in serialization.

Additionally, the model is converted to ONNX format

```
export/rock/saved_model.pb  
export/rock/saved_model.onnx
```

Post-training model quantization

The synthesis script quantizes the weights from provided onnx file internally. if needed to run through evaluation, it can be run with **--generate-dequantized-onnx-file** and it will regenerate a dequantized onnx file that can be used with onnx runtime for evaluation.

Model evaluation

After quantization, the dequantized model can be evaluated and compared with unquantized model (MNIST example) :

```
$ bash evaluate_mnist.sh
```

Alternatively, the user can evaluate all model examples by running bash script:

```
$ bash evaluate_ALL.sh
```

MAX78000 synthesis

To quantize TensorFlow model and synthesize MAX78000 C source code from ONNX file execute the following command (MNIST example) :

```
$ (ai8x-synthesis) ./ai8xsize.py --verbose -L --top-level cnn --test-dir  
tensorflow --prefix tf-mnist --checkpoint-file ../ai8x-  
training/TensorFlow/export/mnist/saved_model.onnx --config-file  
./networks/mnist-chw-ai85-tf.yaml --sample-input ../ai8x-  
training/TensorFlow/export/mnist/sampleddata.npy --device MAX78000 --compact-data  
--mexpress --embedded-code --scale 1.0 --softmax --display-checkpoint $@
```

It requires three input files:

/ai8x-training/TensorFlow/export/mnist/saved_model.onnx - ONNX presentation of TensorFlow model

/ai8x-synthesis/networks/mnist-chw-ai85-tf.yaml - YAML description of the model

/ai8x-training/TensorFlow/export/mnist/sampleddata.npy - Input data sample file

Parameter	Description
--keep-first	If present, it applies same scale factor of weights as specified in --scale to the first layer, otherwise half of the scale factor is applied to the first layer. See Datasets section.
--scale	Scale factor of weight's quantization (default =1)
--generate-dequantized-onnx-file	Generates a dequantized copy of the onnx file to be used for evaluation. See Post-training model quantization section.

Other used parameters are described in section "Network Loader (AI8Xize)" of [1]

Generated C code is stored in **/ai8x-synthesis/tensorflow/tf-mnist/** directory.

To generate MAX78000 C source code for all TensorFlow examples execute following script:

```
$ (ai8x-synthesis) bash gen-tf-demos-max78000.sh
```

Expected shape of input data sample file:

Example	Tensorflow dataset native shape	Synthesis expected shape	commandline option for training to generate expected sample file shape
mnist, fashionmnist (or cases with conv2d as input with 1 channel)	(1,28,28)	same: (1,28,28)	none
cifar10, cifar100, (or cases with conv2d as input with multiple channels like rock)	(1,32,32,3)	channel first: (3,32,32)	--channel-first
kws20 (or cases with conv1d as input)	(1,128,128)	channel first: (128,128)	-- channel-first --swap

References

[1] [ai8x-training/README.md](#)